

Data Warehousing Indexing

Esteban Zimányi

ezimanyi@ulb.ac.be

Slides by Toon Calders

Summary

- How is the data stored?
 - Relational database (ROLAP)
 - Specialized structures (MOLAP)
- How can we speed up computation?
 - Materialized views
 - Indexing structures
 - Partitioning

How Does it Fit In?

- We know what part of the full cube we want to materialize, and how to store it.

We made the problem smaller but did not solve it

- Before partial materialization:

Answer (supplier) from (part, supplier, customer)

- After partial materialization:

Answer (supplier) from (supplier, customer)

How Does it Fit In?

- Not all queries are of the type

```
SELECT D1, ..., Dk, sum(M)
FROM R
GROUP BY D1, ..., Dk
```

How Does it Fit In?

- Example of another type of query:

```
SELECT supplier, year, min(price)
FROM "cube"
WHERE
    product = "toilet paper"
    and (year = 2009 or year = 2010)
GROUP BY supplier, year
```

Indexing Principle

No index




Indexing Principle

No index



INDEX	
Adams, Jesse	60-61
Alvarado, Hiram O.	61
Arnold, Dan and Benina	61-62
Arnold, George and Agatha	62
Arnold, Henry and Cora	18-19
Assembly of God Church (Dilley)	17
Assembly of God Church (Pearsall)	58
Avant, Forrest J.	19
Avant, James Ross	19-20
Avant, Robert F. and Elsie	20
Collins, Clemmons	26
Conover, Benjamin Edward	46
Conover, B. F.	46, 138
Conover, Freddie Marvin	46, 138
Conover, Fred N.	46-47
Conover, George W.	47
Conover, George Washington	47
Conover, Mac D.	47
Conover, Minnie	47
Conover, William O.	47
County, Roosevelt and Lois	69-70
Cowden, George	70-71
Cowley, W. B.	71-72
Cox, Joseph	72



Indexing Principle

- Database Equivalent

No index



Expensive

Full table scan

Index



Inexpensive


index lookup

+ Retrieve data page

Why not Just Use B-Trees?

- The RDBMS work horse!

A B+ Tree



- Make index on (Country,Category,Brand,Chain)

Why not Just Use B-Trees?

B(+)-Trees no longer suffice

- Fixed order between attributes

Index(A,B,C) on R supports:

- Selections on A, on AB, on ABC

Does not support:

- Selections on B, BC, C

We need exponentially many B-trees to support all possible selections

- Attributes are spread over different tables

Summary

- How can we speed up computation?
 - Materialized Views
 - Indexing structures
 - Bitmap index
 - Join index
 - Bitmap-join index
 - Partitioning

Bitmap Index: Definition

- Indexing structure for
 - One attribute $R.A$ for one relation R
 - Optimizing queries making **Boolean combinations of selections** on indexed attributes
- For every value v in the active domain of A store a bitmap
 - Length of the bitmap = size of the relation R
 - Bitmap has value 1 on position k if the k th tuple of R has value v in attribute A

Bitmap Index: Example

Product	Country	Sales
TV	Ireland	20
TV	France	126
HiFi	Germany	56
PC	Ireland	23
TV	France	138
PC	Germany	48

- Index for Country:

Ireland	100100
France	010010
Germany	001001

Bitmap Index: Example

Product	Country	Sales
TV	Ireland	20
TV	France	126
HiFi	Germany	56
PC	Ireland	23
TV	France	138
PC	Germany	48

- Index for Country:

Ireland	100100
France	010010
Germany	001001

Index for Product

TV	110010
HiFi	001000
PC	000101

Bitmap Index: Example

- Index for Country:

Ireland	100100
France	010010
Germany	001001

- Index for Product

TV	110010
HiFi	001000
PC	000101

```
SELECT sum(Sales)
```

```
FROM PCS
```

```
WHERE (Country = Ireland or Country = France)  
      and not(Product = TV)
```

Access only tuples corresponding to a 1 in the bitmap:

$$(\boxed{100100} \mid \boxed{010010}) \& ! \boxed{110010} = \boxed{000100}$$

Bitmap Index: Space

- Size of bitmaps can be reduced

E.g., **run-length encoding** (RLE):

1111000111100000001111000 is encoded as

4x1;3x0;4x1;7x0;4x1;3x0

kati ota 1 kati ota 0


- Reduces storage requirements significantly
 - Logical operations can **work directly on RLE**
-
- BTree maps values to the (variable length) bitmaps

Working Directly on RLE

- Brussels: 0000000011000010000000
7x0, 2x1, 4x0, 1x1, 6x0
- TV: 00010001110010000000
3x0, 1x1, 3x0, 3x1, 2x0, 1x1, 7x0
- Orange: 00000111110101110000
5x0, 5x1, 1x0, 1x1, 1x0, 3x1, 4x0
- !(orange) & (Brussels | TV):
!orange: 5x1, 5x0, 1x1, 1x0, 1x1, 3x0, 4x1
Brussels | TV: 3x0, 1x1, 3x0, 3x1, 2x0, 2x1, 6x0
!o & B|T: 3x0, 1x1, 8x0, 1x1, 7x0

Bitmap Index

bitmap along with btree



Leafs contain the bitmaps

Variants of Bitmap Index

Disjoint ranges
use case - age

Overlapping ranges
use case: for temperature
because we might have queries most frequently asked as get all
temp>10

which one is better depends on type of query we will be doing

- Numerical attributes

- Discretize into intervals

E.g., temperature into $] -10, -5]$, $] -5, 0]$, $] 0, 5]$, $] 5, 10]$, ...

- Make one bitmap for each bin

- Combine bitmaps at query time

temperature > 5 \rightarrow $b(]0, 5]) \mid b(]5, 10]) \mid \dots$

- Other option: make overlapping bitmaps

- Bitmap for all tuples with $t > -10$, $t > -5$, $t > 0$, $t > 5$, ...

temperature between 0 and 5: $b(t > 0) \& \neg b(t > 5)$

Variants of Bitmap Index

- Bit-sliced index for nominal attribute with many values:
 - Encode values as integers
 - Make bitmap for bit 0, for bit 1, for bit 2, ...

indexing string values

Customer	City
C1	Antwerp
C2	Brussels
C3	Eindhoven
C4	Eindhoven

Variants of Bitmap Index

- Bit-sliced index for nominal attribute with many values:
string represented in the form of integer
and somewhere else will be having mapping
 - Encode values as integers
 - Make bitmap for bit 0, for bit 1, for bit 2, ...

Customer	City
C1	0000
C2	0001
C3	1011
C4	1011

Bit 1	Bit 2	Bit 3	Bit 4
0	0	0	0
0	0	0	1
1	0	1	1
1	0	1	1

Variants of Bitmap Index

- Bit-sliced index for nominal attribute with many values:
 - Encode values as integers
 - Make bitmap for bit 0, for bit 1, for bit 2, ...

Customer	City
C1	0000
C2	0001
C3	1011
C4	1011

Bit 1	Bit 2	Bit 3	Bit 4
0	0	0	0
0	0	0	1
1	0	1	1
1	0	1	1

Antwerp or Brussels



$(!b1 \ \& \ !b2 \ \& \ !b3 \ \& \ !b4) \mid (!b1 \ \& \ !b2 \ \& \ !b3 \ \& \ b4)$
 $(!b1 \ \& \ !b2 \ \& \ !b3)$

Bitmap Index: Pros and Cons

if the value has updated find that row, and update value accordingly, which is hard due to compression
more useful for read only cases

- Bitmap Indices are **hard to maintain**
 - Due to compression
 - Usually **rebuilt after bulk update**
- Particularly useful when there are **multiple** bitmap indices
- *For **low cardinality** attributes*
 - Bitmaps lose their edge when cardinality is high

not useful for price
if we need to do, can create ranges and apply bin concept

Index Combination: Bitmap Filtering

- Not all systems support bitmap indices (e.g., SQL Server)
 - Yet, has a different way to combine indices
 - bitmap filtering

Index on A → list of rowIDs

Index on B → list of rowIDs

Intersect lists

Index Combination: Bitmap Filtering

- Not all systems support bitmap indices (e.g., SQL Server)
 - Yet, has a different way to combine indices
→ bitmap filtering

Index on A → list of rowIDs → bitmap

Index on B → list of rowIDs → bitmap

bitmap arithmetics

(Vertical) Projection Index

- Almost the same
 - Instead of storing bitmaps: project on columns
 - Keep all projections in the same order

between two refresh these values are kept in same order

Customer	Age	City
C1	30	Antwerp
C2	34	Brussels
C3	39	Eindhoven
C4	65	Eindhoven



Customer	City	Age
C1	Antwerp	30
C2	Brussels	34
C3	Eindhoven	39
C4	Eindhoven	65

Added to the database

(Vertical) Projection Index

- Advantage:
 - better data locality if only few attributes are needed
needed info are brought from disk to memory
 - No separate, hard to maintain index needed
- Column-store databases have proven to be highly performant for managing large datasets

Remark: Bitmap and Projection Index

- Mapping bitmaps to tuples not always straightforward
 - For instance, where to locate 5th tuple?
- If tuples are stored consecutively, and have equal length: $\text{offset} + (\text{nr} - 1) * \text{length}$
- Otherwise: next to bitmaps array with physical addresses
 - entry i in array holds physical location of i th tuple
 - Only one such array per table needed

Summary

- How can we speed up computation?
 - Materialized Views
 - Indexing structures
 - Bitmap index / Projection index
 - Join index
 - Bitmap-join index
 - Indexing fact and dimension tables
 - Partitioning

Join Index

- Traditional indexes:
 - value of R.A \rightarrow rows/RIDs in R
- Join indices:
 - Value of R.A \rightarrow RIDs in S
- Data warehouse:
 - Attribute in dimension table \rightarrow rows in fact table
 - Join indexes can span multiple dimensions


Join Index: Example

Sales

Date	pID	Client
10/5/12	1	Jack
10/5/12	1	Pete
13/5/12	3	John
14/5/12	2	Mary

Products

pID	pName	Category	Price
1	Jacket	Non-food	10
2	Bread	Food	2.1
3	Beer	Food	1.5
4	Paper	Non-food	1.2



Date	pID	Client	pName	Category	Price
10/5/12	1	Jack	Jacket	Non-food	10
10/5/12	1	Pete	Jacket	Non-food	10
13/5/12	3	John	Beer	Food	1.5
14/5/12	2	Mary	Bread	Food	2.1


Join Index: Example

Sales

Date	pID	Client
10/5/12	1	Jack
10/5/12	1	Pete
13/5/12	3	John
14/5/12	2	Mary

Products

pID	pName	Category	Price
1	Jacket	Non-food	10
2	Bread	Food	2.1
3	Beer	Food	1.5
4	Paper	Non-food	1.2



SP_category_jidx

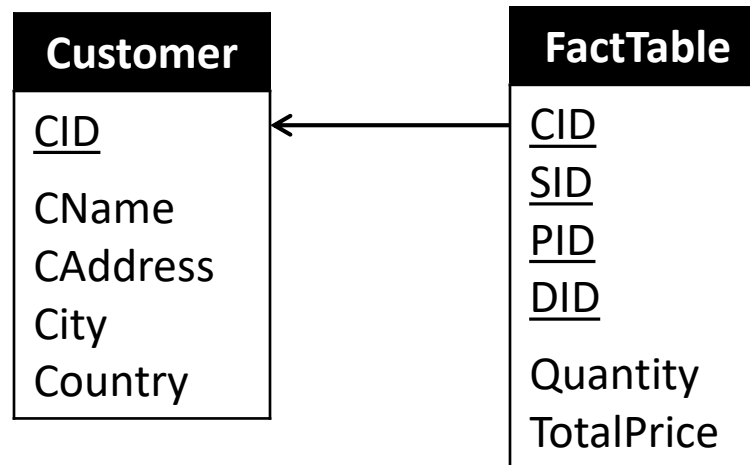
Category	RID
Non-food	r1,r2
Food	r3,r4

cat materialized

Date	pID	Client	pName	Category	Price
10/5/12	1	Jack	Jacket	Non-food	10
10/5/12	1	Pete	Jacket	Non-food	10
13/5/12	3	John	Beer	Food	1.5
14/5/12	2	Mary	Bread	Food	2.1

Join Index: Data Warehouse

- Join index can index tuples in the fact table based on an attribute in a dimension



- E.g., Index tuples in the fact table for the attribute Country of Customer

Join Index: Variant

- If index attribute is primary key of one table:
directly store RIDs into the table itself
 - Avoids lookup in index when conditions on product.

instead of repeating non food two time
we can add lookup

Products

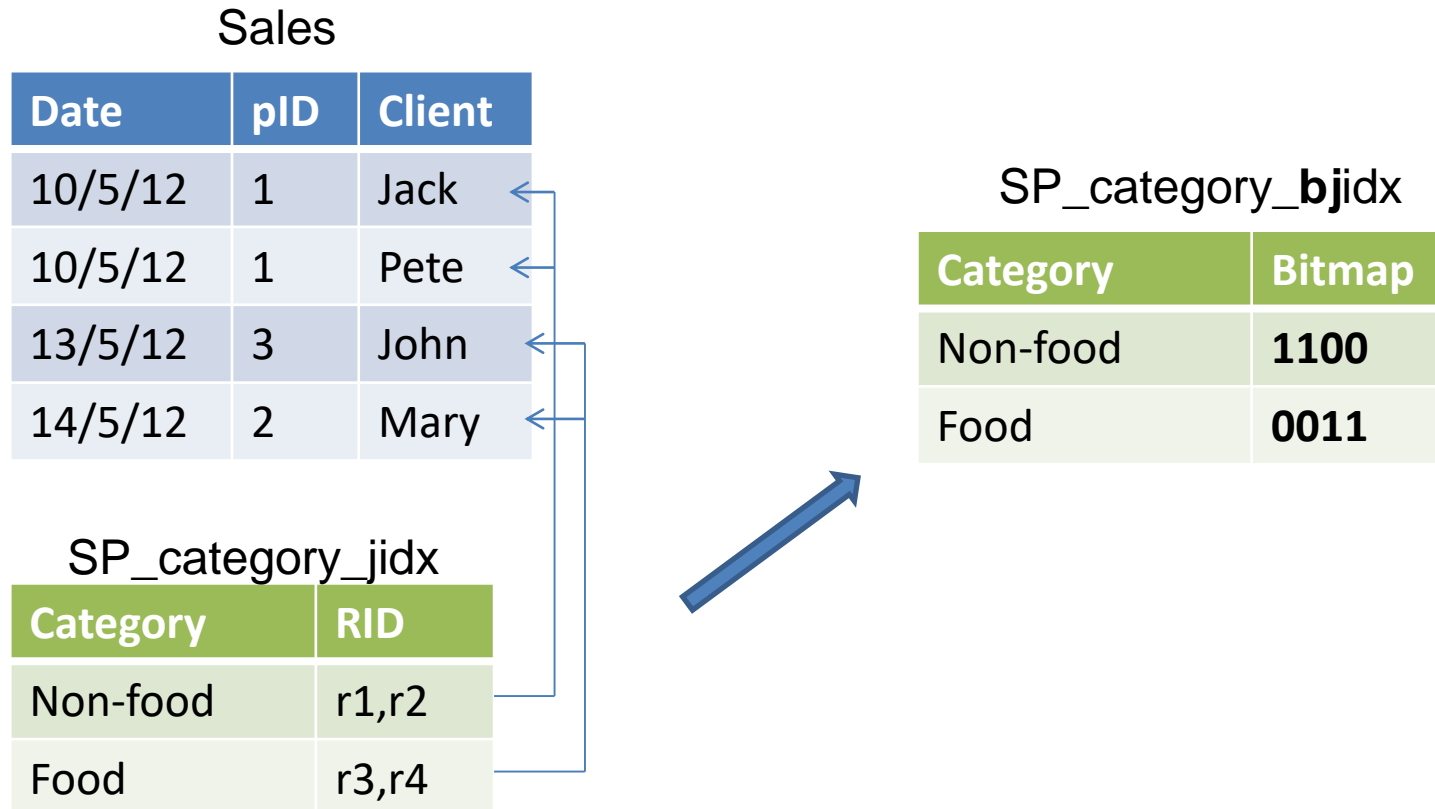
pID	pName	Category	Price	RIDSales
1	Jacket	Non-food	10	r1, r2
2	Bread	Food	2.1	r4
3	Beer	Food	1.5	r3
4	Paper	Non-food	1.2	

Sales

Date	pID	Client
10/5/12	1	Jack
10/5/12	1	Pete
13/5/12	3	John
14/5/12	2	Mary

Bitmap-join index

- Logical combination of bitmap index and join index



Bitmap-Join Index: Example

Date	pID	Client
10/5/12	1	Jack
10/5/12	1	Pete
13/5/12	3	John
14/5/12	2	Mary

SP_category_bjidx

Category	Bitmap
Non-food	1100
Food	0011

SC_city_bjidx

City	Bitmap
Brussels	1001
Eindhoven	0110

```
SELECT date
FROM Sales S JOIN Product
      P JOIN Customer C
WHERE
      P.Category = "Food" and
      C.City = "Brussels";
```

Bitmap-Join Index: Example

Date	pID	Client
10/5/12	1	Jack
10/5/12	1	Pete
13/5/12	3	John
<u>14/5/12</u>	2	Mary

SP_category_bjidx

Category	Bitmap
Non-food	1100
Food	0011

SC_city_bjidx

City	Bitmap
Brussels	1001
Eindhoven	0110

SELECT date

FROM Sales S JOIN Product
P JOIN Customer C

WHERE

P.Category = "Food" and
C.City = "Brussels";

0011 & 1001 → 0001

Indices in Practice

- Several commercial products implement bitmap, join & bitmap-join index
 - E.g., Oracle (EE) offers a compressed bitmap index

```
CREATE BITMAP INDEX cust_sales_bji
ON sales(customers.city)
FROM sales, customers
WHERE sales.client = customers.name;
```
- Some products build bitmaps *on the fly*
 - E.g., SQLServer 2008 → *bitmap filtering*

Summary

- How can we speed up computation?
 - Materialized Views
 - Indexing structures
 - Bitmap index / Projection index
 - Join index
 - Bitmap-join index
 - Indexing fact and dimension tables
 - Partitioning

Indexing Fact and Dimension Tables

- Type of index depends on type of attributes

higher the age criteria less chance of getting selected in university records

		Distinct values	
		Few	Many
Query selectivity	Low	Bitmap	Compressed bitmap Projection index B+-tree
	High	B+-Tree Bitmap	B+-tree

- Optimal set of indices depends on workload
 - attribute often used for slicing → bitmap index
 - Key attribute → B+-tree

Indexing Dimension Tables

- Useful indices:
 - Dimensional attribute that is often used for slicing
 - Low cardinality → bitmap/bitmap-join index
 - High cardinality → B+-tree index
 - B+-tree index on surrogate key
(necessary for joins with the fact table!)
 - Foreign keys (snowflake schema)

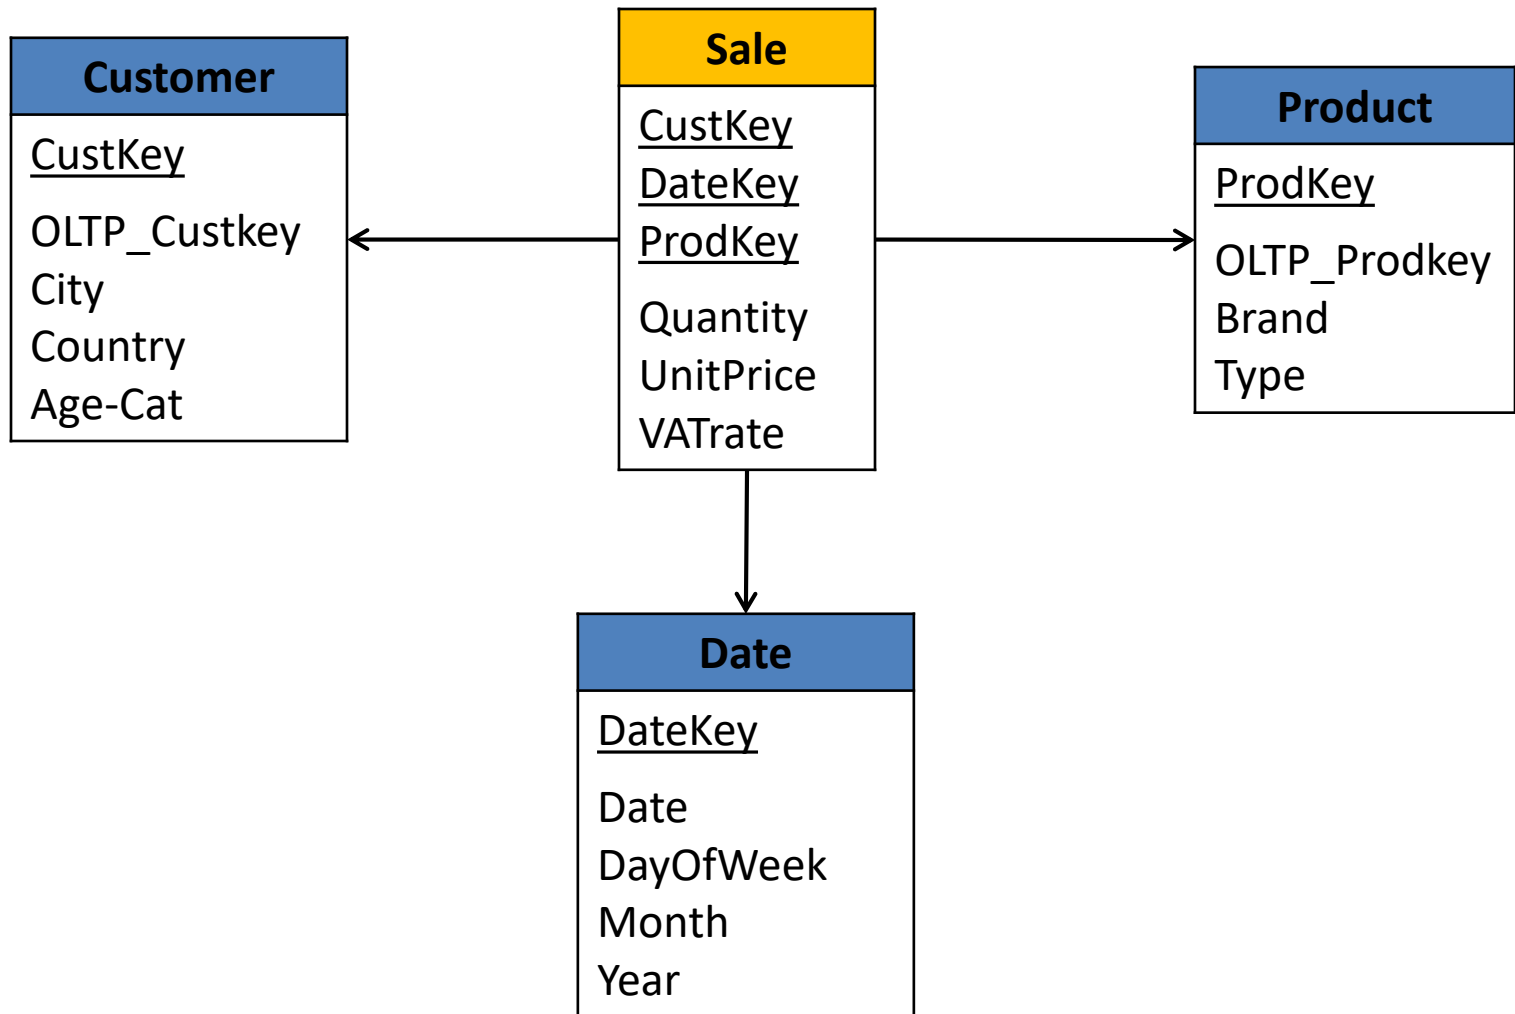
**

Ex: Client City State Country
Snowflake schema - B-tree Index or bitmap join index
Star schema - bitmap index or bitmap join index

Indexing the Fact Table

- Index(es) on the foreign keys
(for joining with dimension tables)
 - Index on attributes (A,B,C) helps queries that join on A, AB, ABC, but not on the other attribute sets
 - Often useful to have multiple indices for different orders/for foreign keys separately

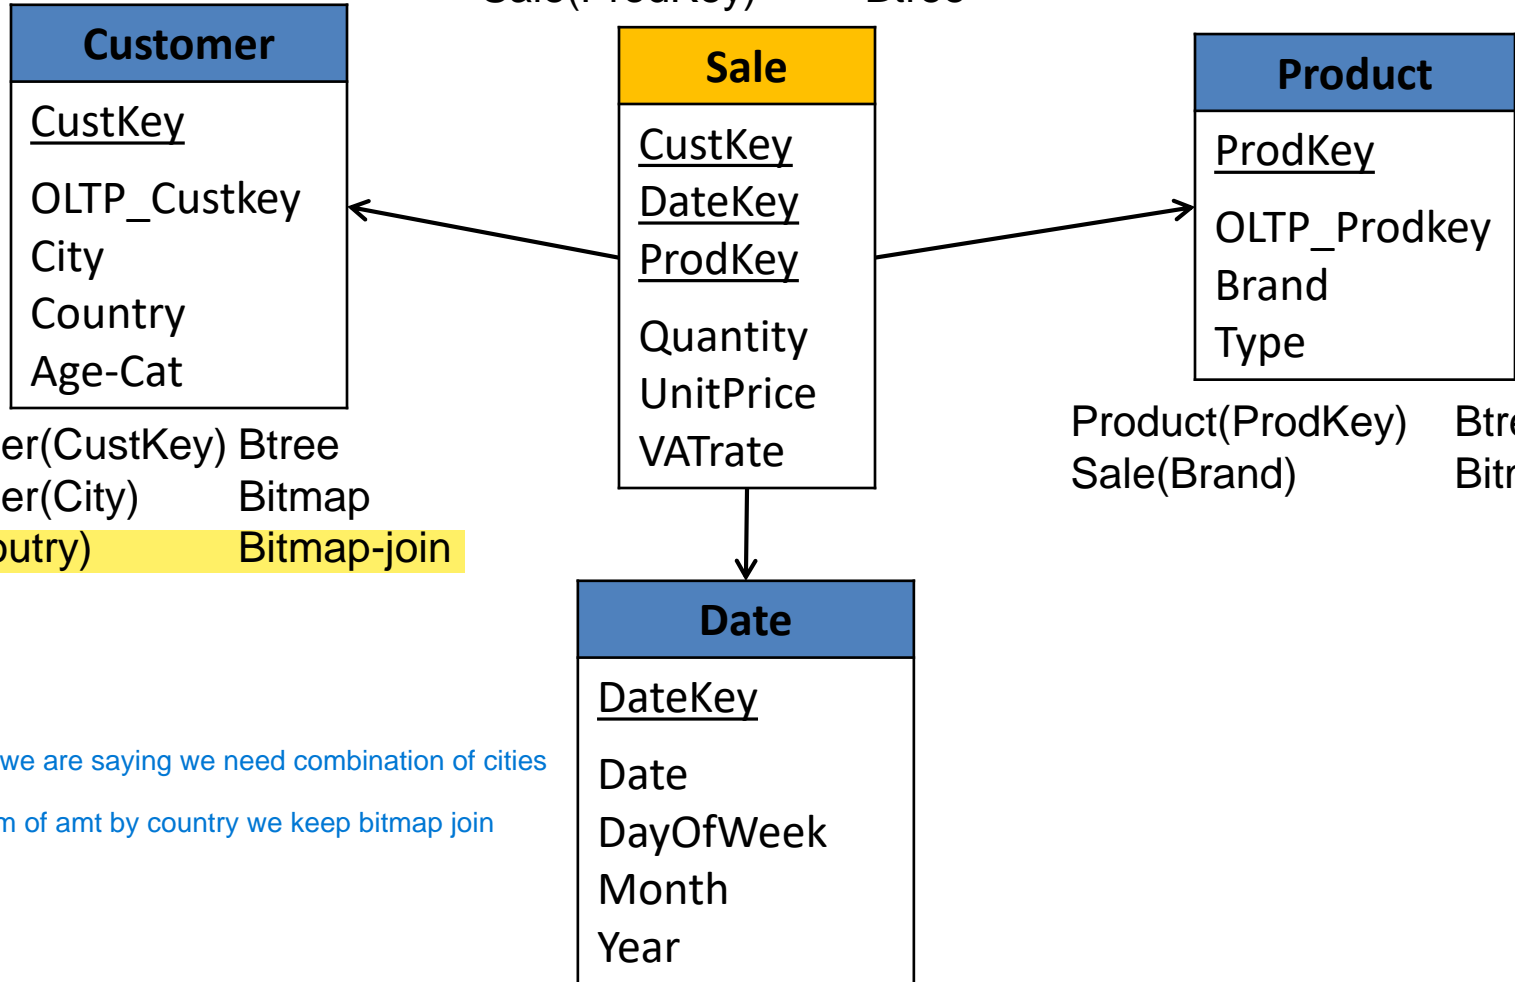
Example



Example

we don't keep index for product name because we already have prodkey and have high variance as well

Sale(CustKey) Btree
 Sale(DateKey) Btree
 Sale(ProdKey) Btree



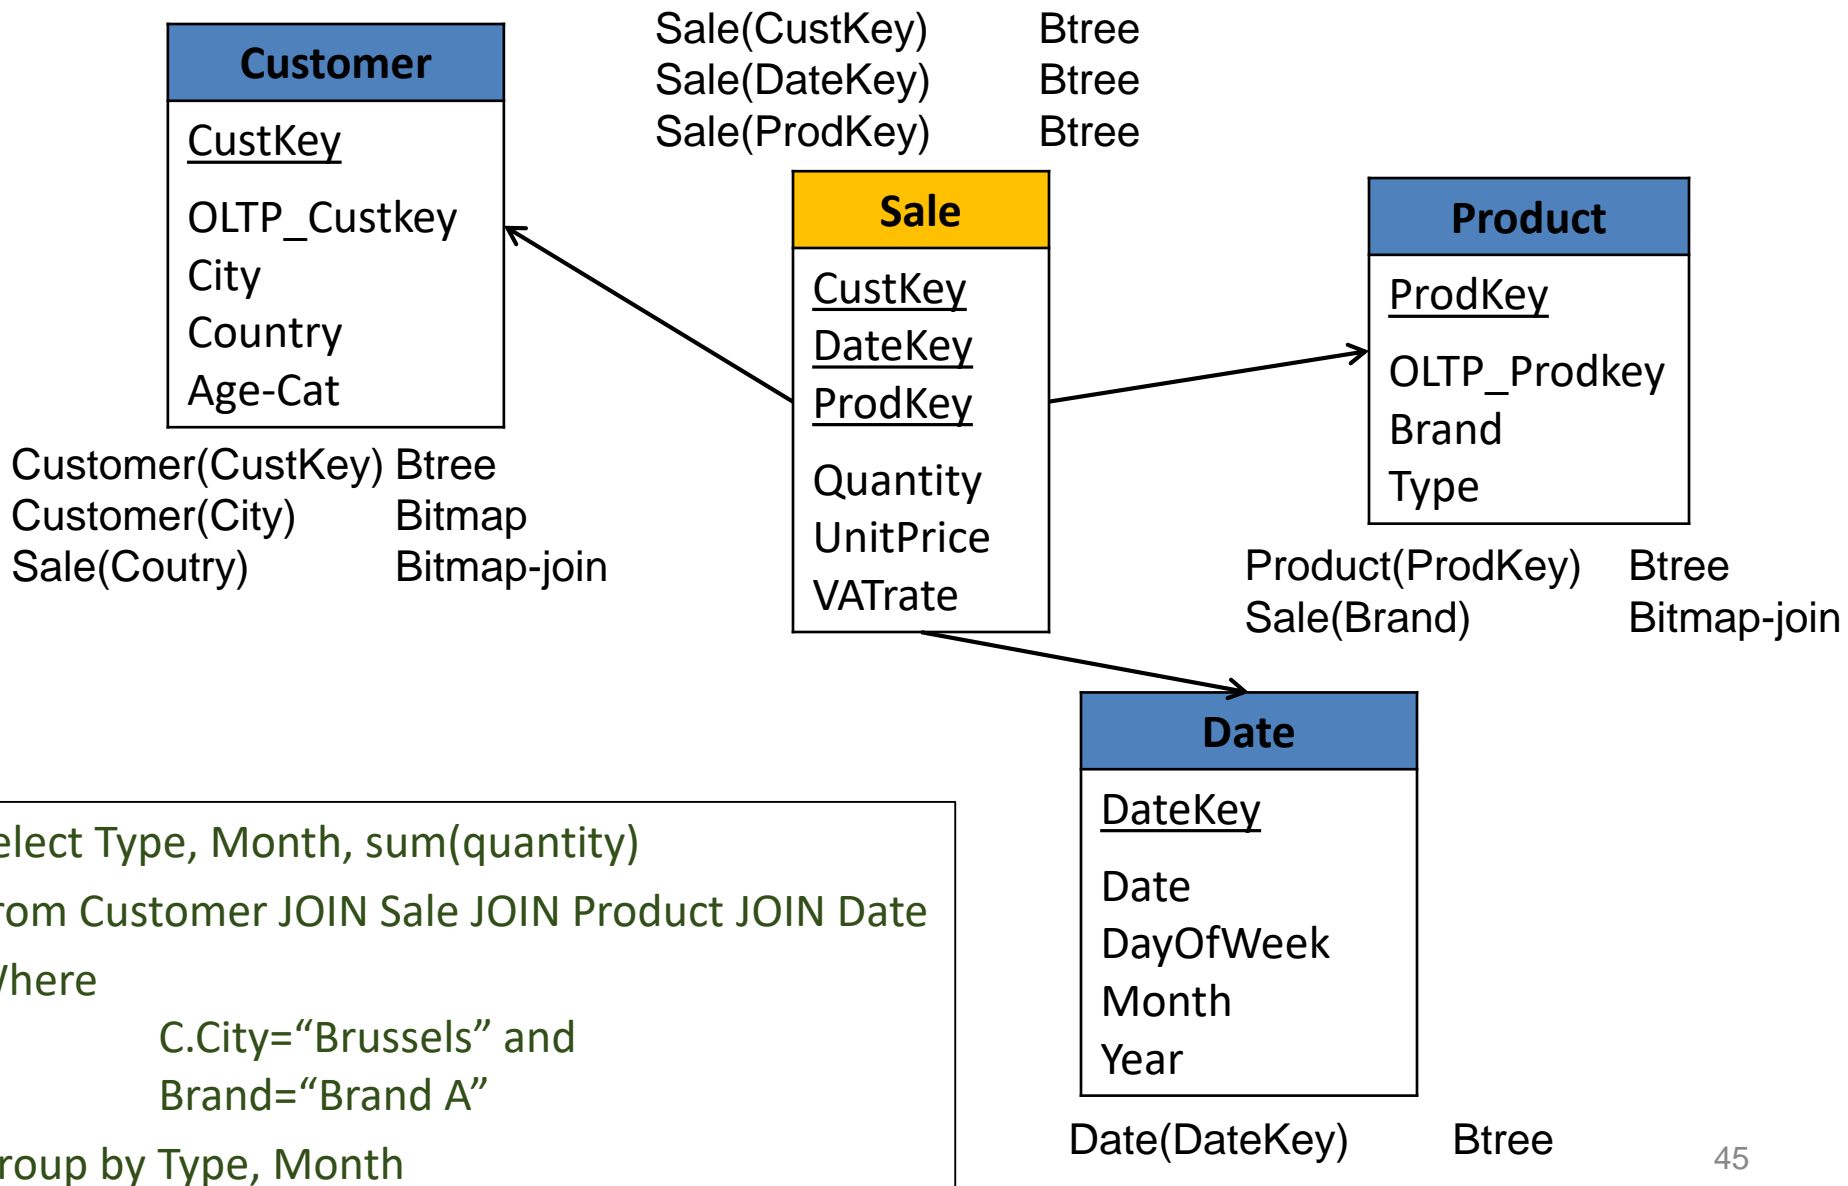
Customer(CustKey) Btree
 Customer(City) Bitmap
 Sale(Country) Bitmap-join

Product(ProdKey) Btree
 Sale(Brand) Bitmap-join

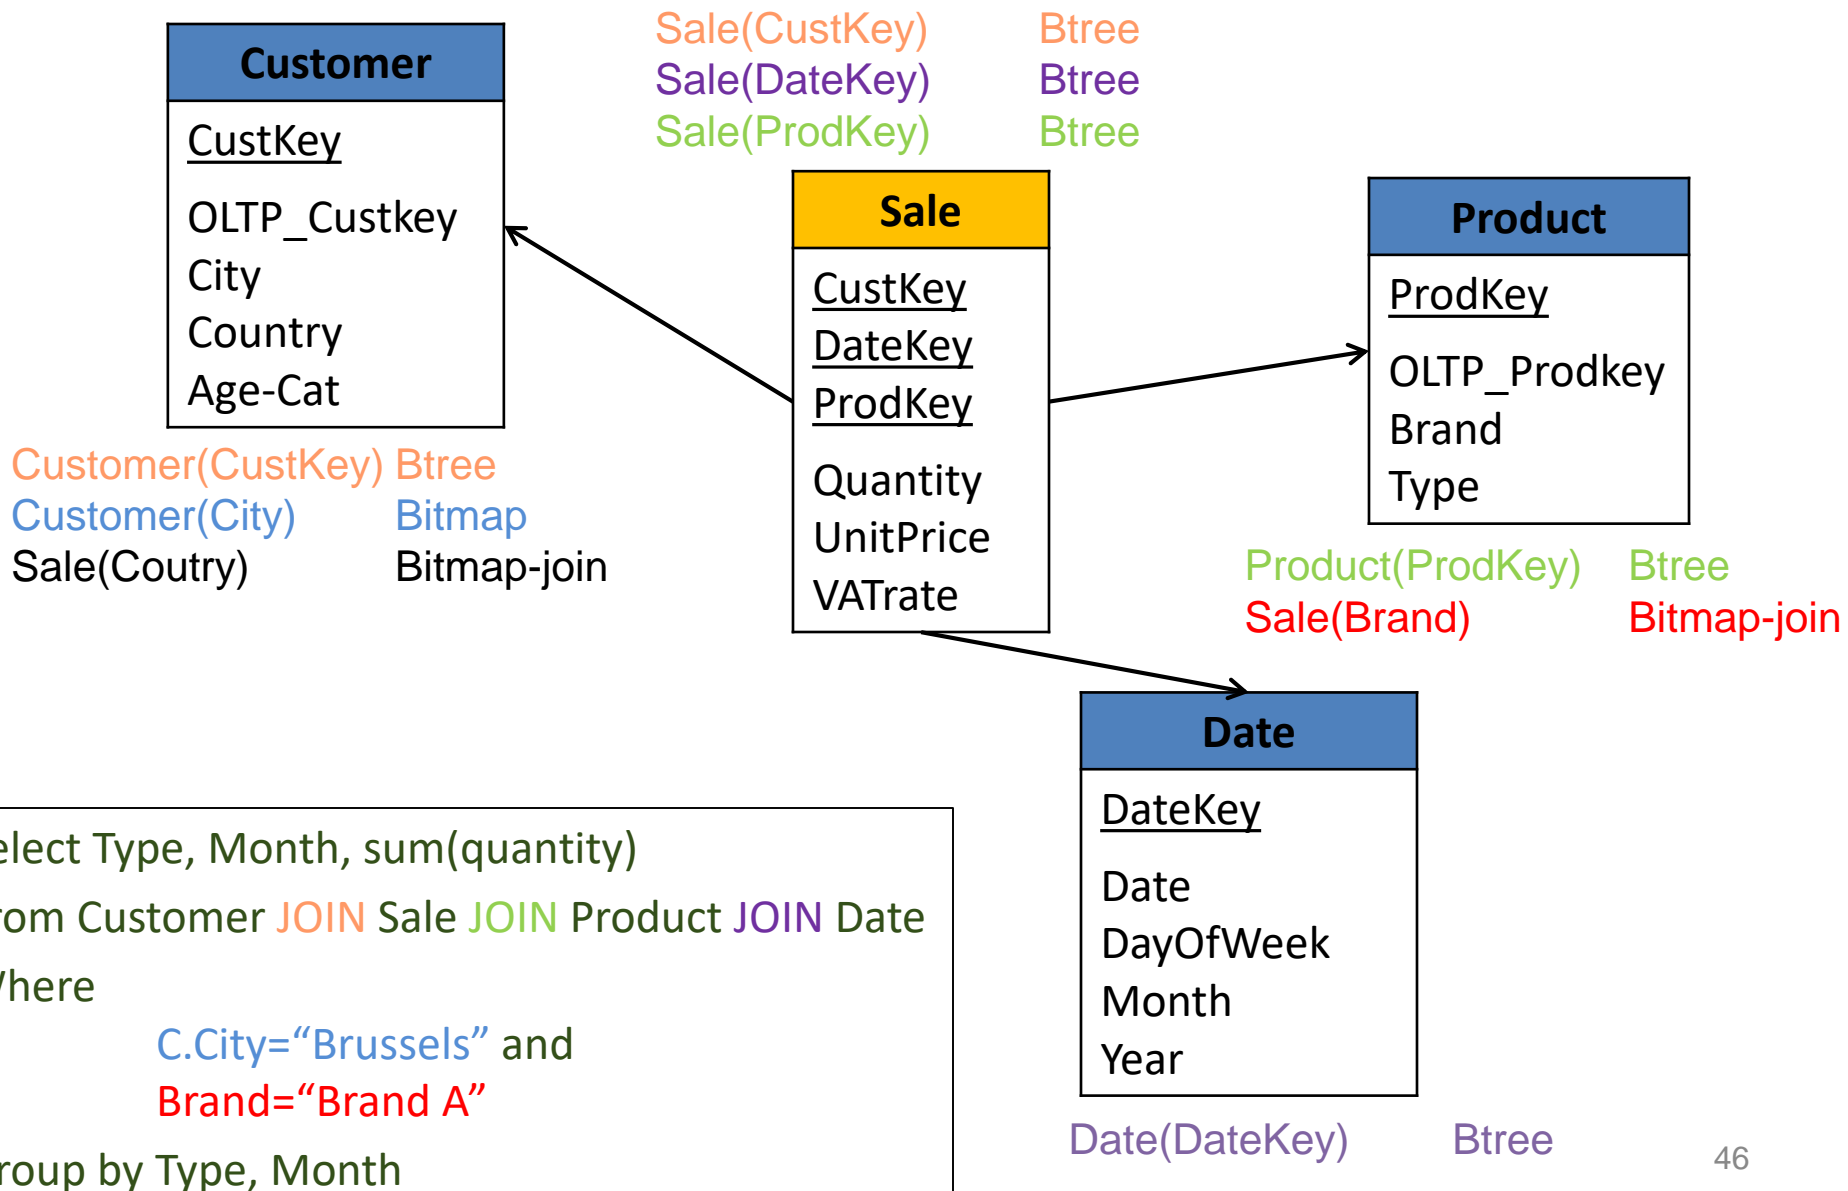
bitmap halda we are saying we need combination of cities
 if we want sum of amt by country we keep bitmap join

Date(DateKey) Btree

Example



Example



Example

- Different options:
 - Use **bitmap-join index** to directly select facts about “Brand A”
 - Aggregate by Type, Cust, Date (reduce # tuples in join)
 - Use Btrees on primary keys in **Customer** and **Date** table to join in these tables (index nested-loop)
 - Filter tuples of join with “City=Brussels”
 - Group by Type, Month

```
Select Type, Month, sum(quantity)
From Customer JOIN Sale JOIN Product JOIN Date
Where
    C.City="Brussels" and
    Brand="Brand A"
Group by Type, Month
```

Example

- Different options:
 - Use **bitmap-join index** to directly select facts about “Brand A”
 - Aggregate by Type, Cust, Date (reduce # tuples in join)
 - Filter Customer on “City=Brussels” using **bitmap idx**
 - Join filtered Customer with Fact table (nested loop)
 - Use **Btree on primary keys in Date** table to join with fact table (index nested-loop)
 - Group by Type, Month

```
Select Type, Month, sum(quantity)
From Customer JOIN Sale JOIN Product JOIN Date
Where
    C.City="Brussels" and
    Brand="Brand A"
Group by Type, Month
```


Example

- Different options:
 - Filter Customer on “City=Brussels” using **bitmap idx**
 - Join Customer and fact table using **Btree on foreign key CustKey** in fact table (index nested-loop)
 - Use Btree on primary keys in **Date** and **Product** table to join (index nested-loop)
 - Filter on Brand=“Brand A”
 - Group by Type, Month

```
Select Type, Month, sum(quantity)
From Customer JOIN Sale JOIN Product JOIN Date
Where
    C.City="Brussels" and
    Brand="Brand A"
Group by Type, Month
```

Example

- Based on an estimate of the cost, optimizer will select the cheapest plan
 - Very hard to predict
 - Depends on database statistics
- Based on usage and database statistics index selection should be revised regularly
 - Remove indices that are not used
 - Add indices to speed up slow queries

Summary

- How can we speed up computation?
 - Materialized Views
 - Indexing structures
 - Bitmap index / Projection index
 - Join index
 - Bitmap-join index
 - Indexing fact and dimension tables
 - Partitioning

Partitioning

- Separate database/tables/indices over different partitions
 - Horizontal partitioning: every partition holds a subset of the *tuples*
E.g., partition fact table by *month*
 - Vertical partitioning: every partition holds a subset of the *attributes*

Partitioning

- Advantages of horizontal partitioning
 - Easier for data warehouse refresh
 - No need to rebuild index for the whole table
 - Ease of maintenance; e.g., removing an outdated partition
- Disadvantage:
 - Overhead & reduces efficiency of indexing, especially if query spans many partitions
 - optimal partitioning depends on workload

Conclusion

- Bitmap index, join-index, bitmap-join index:
 - Speedup selection queries with arbitrary Boolean combinations of indexed attributes
 - Very interesting for ad-hoc analytical queries
 - Not easy to update
 - Not suitable for operational databases: inserts and deletes
 - Typically these indices are completely rebuild after *bulk inserts*

Typical for Data Warehouses; less suitable for OLTP